

Phase 3 Presentation

Reinforcement Learning for Adaptive Numerical Integration

Team KVS

Kavya Bhalodi (2023101031)

Varun Gupta (2023101108)

Shivam Gupta (2023101062)

International Institute of Information Technology, Hyderabad

May 8, 2025

Motivation

- Traditional adaptive integration relies on fixed heuristics.
- Challenging for functions with non-smooth features or varying complexity.
- Goal: Use RL to dynamically adapt evaluation effort based on function behavior.
- Potential for better accuracy-effort tradeoff.

Previous Work: 1D RL-Based Integration

- Developed an RL-based adaptive integration model for 1D functions.
- Used 5-point Gauss-Legendre Quadrature for numerical estimation.
- Richardson Extrapolation was employed to estimate integration error.
- The PPO agent learned to split intervals based on:
 - Function slope, curvature, variance, and extrapolated error.
- Demonstrated improved accuracy-effort trade-off over classical methods.

Previous Work: 1D RL-Based Integration

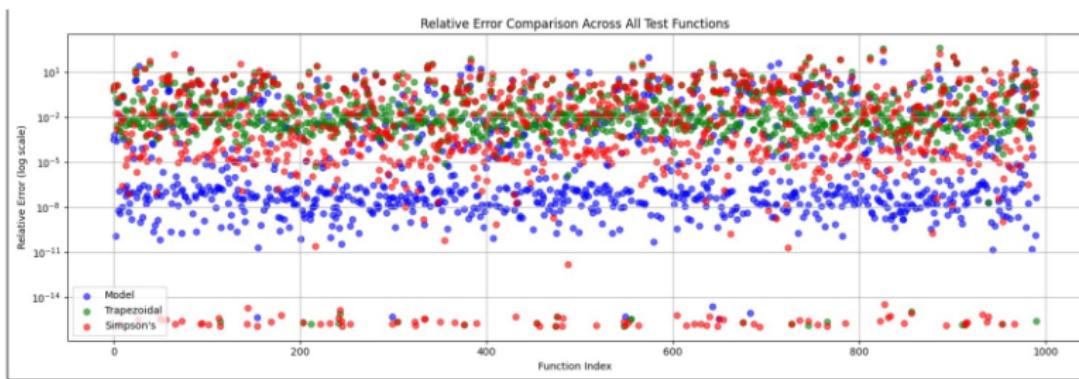


Figure: Relative error across 1000 test functions. Our model achieves significantly lower error than Trapezoidal and Simpson's methods.

Figure: Enter Caption

Problem Setup: 2D Integration

- Evaluate integrals of the form $\iint f(x, y) dx dy$ using adaptive strategies.
- Agent learns to refine rectangular subregions based on function complexity.
- State includes:
 - Region size and geometric bounds
 - Function behavior: curvature, oscillation, skew, and kurtosis
 - Estimated error
- Action: Select region, split axis, and strategy (e.g., midpoint vs. adaptive point)
- Reward: Based on global and local error reduction per evaluation cost

Our Approach

- We use a **Reinforcement Learning (RL)** agent to adaptively refine sub-intervals during 2D numerical integration.
- The function domain is divided into rectangular cells.
- At each step, the agent decides **which interval to refine** based on:
 - Local function behavior (e.g., curvature, error)
 - Features extracted from sampled points
- Rewards are designed to encourage high accuracy with fewer evaluations.
- Integration methods (Adaptive Simpson, Monte Carlo, Gaussian, PMEF) are plugged in depending on scenario.

Methodology

- Use PPO (Proximal Policy Optimization) to train an RL agent.
- Agent traverses a function tree, choosing where to split.
- Chebyshev points, which are specific points chosen to optimize accuracy, are used to evaluate the function in each sub-interval.
- Policy updated via policy gradients using real-valued reward.

Numerical Methods Used

Quadrature Techniques

- 5-point Gauss-Legendre Quadrature (2D)
- Monte Carlo Integration
- PFEM (Particle Finite Element Method)
- Adaptive Simpson Method

Error Estimation

- **Richardson Extrapolation:** Compares coarse vs. fine grid Gauss estimates to compute local error.
- **Statistical Measures:** Uses standard deviation of neighboring particle values and Monte Carlo variance to assess error dynamically.

Adaptive Strategy

- Environment chooses between quadrature, Monte Carlo, or PFEM based on local function behavior (oscillation, smoothness, curvature).
- Error-driven region splitting: RL agent selects where and how to split domains using function statistics and estimated errors.

Numerical Integration Methods

- **Adaptive Simpson Method:** Approximates the integral over a grid by fitting a quadratic surface; accurate for smooth functions.
- **Monte Carlo Integration:** Uses random sampling over the domain; useful for complex or irregular regions.
- **Gaussian–Legendre Quadrature (5-point):** Applies Gaussian quadrature in both dimensions with 5 nodes; highly accurate for smooth, bounded regions.
- **PFEM (Particle Finite Element Method):** Deployed for sharp features or high-gradient areas. Uses particles with adaptive refinement based on local error.

These methods are selected dynamically or fixed depending on the experiment setup and function characteristics.

PFEM-Based Integration: Overview

Particle Finite Element Method (PFEM)

- Used for numerically integrating regions with:
 - High oscillation, curvature, or gradient
- Initializes particles using jittered grid sampling over the region.
- Each particle maintains:
 - Position (x, y)
 - Function value $f(x, y)$
 - Local error estimate based on neighbors

Adaptive Particle Distribution

- **Low-error** particles are pruned.
- **High-error** particles trigger localized refinement.
- New particles added around high-error zones using random angular offsets.

PFEM-Based Integration: Computation

Integration Estimate

- Region area $A = (x_1 - x_0)(y_1 - y_0)$
- Final integral computed as a weighted average:

$$I \approx A \cdot \frac{\sum_i f_i w_i}{\sum_i w_i}$$

- Function values f_i are either computed or cached.

Error Estimation

- Neighbors found for each particle within a fixed distance.
- Local error estimated as:

$$\varepsilon_i = \text{std}(\{f_j : j \in \text{neighbors}(i)\})$$

- Global error = mean of all ε_i

Role in RL Environment

- Triggered selectively for challenging regions.
- Balances accuracy with particle computation cost.

PFEM Accuracy Across Challenging Functions

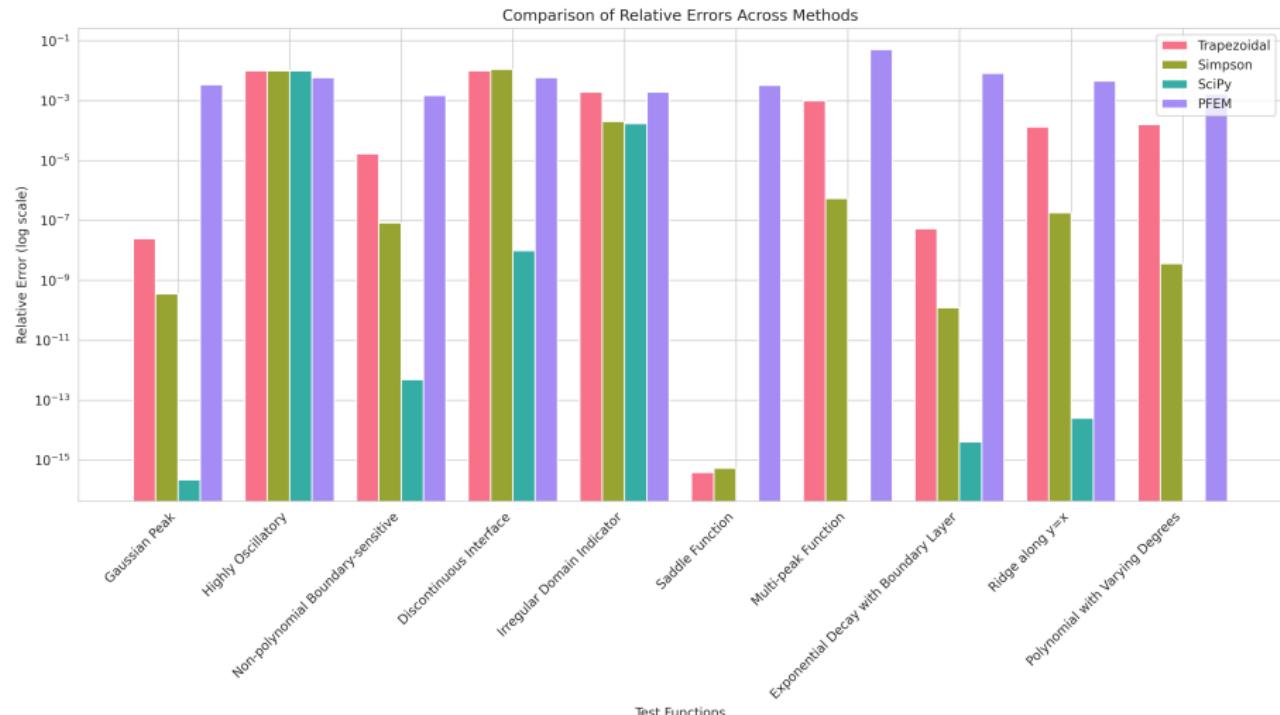


Figure: PFEM achieves significantly lower relative error for highly oscillatory and discontinuous functions, outperforming classical methods in complex regions.

Numerical Integration Method: Gauss-Legendre (2D)

Gauss-Legendre Quadrature (5-point):

- Approximates the integral of $f(x, y)$ over a 2D rectangle $[x_0, x_1] \times [y_0, y_1]$ using tensor product:

$$\iint f(x, y) dx dy \approx \sum_{i=1}^5 \sum_{j=1}^5 w_i w_j \cdot f(x_i, y_j)$$

- x_i, y_j are transformed Gauss points in the 2D domain.
- w_i, w_j : weights corresponding to Gauss points.

Properties:

- High-order accuracy on smooth surfaces.
- Used for estimating integrals and for Richardson extrapolation in 2D.

Numerical Integration Method: Monte Carlo (2D)

Monte Carlo Integration:

- Approximates the integral over region $[x_0, x_1] \times [y_0, y_1]$ using N random samples:

$$\iint f(x, y) dx dy \approx A \cdot \frac{1}{N} \sum_{i=1}^N f(x_i, y_i)$$

- (x_i, y_i) sampled uniformly within the region.
- A : Area of the region.

Usage in Environment:

- Activated for regions with high oscillation, skew, or kurtosis.
- Error estimated using standard deviation of sampled values.
- Simple and robust in noisy or irregular regions.

Numerical Integration Method: Simpson's Rule (2D)

Simpson's Rule (Tensor Product Form):

- Used over structured grids, approximating:

$$\iint f(x, y) dx dy \approx \frac{h_x h_y}{9} (f_{00} + 4f_{01} + f_{02} + 4f_{10} + 16f_{11} + \dots)$$

- Combines Simpson's 1D rule in x and y directions.

Properties:

- Moderate-order accuracy with regular spacing.
- Included for comparison in benchmarking; not used by RL agent.

Error Estimation: Richardson Extrapolation

Richardson Extrapolation:

- Estimates error by comparing coarse and fine Gauss-Legendre approximations:

$$\text{Error} \approx \frac{|I_{\text{fine}} - I_{\text{coarse}}|}{2^n - 1}$$

- n : Order of the integration method.

Usage:

- Fine estimate from applying Gauss-Legendre on split subregions.
- Coarse estimate from applying it once over the full region.

Strengths:

- Reliable for smooth functions.
- Helps guide the RL agent's decision to split or not.

Error Estimation: Monte Carlo Variance

Statistical Error Estimation:

- Error estimated using standard deviation of random samples:

$$\text{Error} \approx A \cdot \frac{\sigma}{\sqrt{N}}$$

- A : Region area, σ : standard deviation, N : number of samples.

Usage:

- Applied in regions where function is non-smooth or highly oscillatory.
- Error passed to RL agent for action selection.

Benefits:

- Robust in irregular or noisy domains.
- Naturally adapts with sampling density.

Error Estimation: PFEM Local Error

PFEM Particle-Based Estimation:

- Each particle computes local error from its neighbors:

$$\varepsilon_i = \text{std}(\{f_j\}_{j \in \mathcal{N}(i)})$$

- Neighborhood $\mathcal{N}(i)$ defined by distance threshold.

Total Error:

$$\text{Global Error} = \frac{1}{N} \sum_{i=1}^N \varepsilon_i$$

Strengths:

- Sensitive to local variations and sharp transitions.
- Helps selectively refine particle density in difficult regions.

RL Model for 2D Adaptive Integration

Objective: Learn to adaptively split 2D regions to minimize function evaluations while maintaining high integration accuracy.

Formulation as a Markov Decision Process (MDP):

- **States:** Capture geometric and statistical features of each rectangular region:
 - Region bounds, area, function center value
 - Function behavior: oscillation, smoothness, skewness, kurtosis
 - Error estimates from Richardson extrapolation or PFEM
- **Actions:**
 - Select region index (normalized)
 - Choose split axis (x or y) and split ratio
 - Strategy type (e.g., midpoint, adaptive sampling)
- **Rewards:** Encourage global and local error reduction per evaluation.
Penalize inefficient splits or poor region selection.

State Representation: 2D Integration

Each state corresponds to a rectangular subregion and is described using:

- **Geometric Features:** Region boundaries, width, height, and area.
- **Function Behavior:**
 - **Oscillation:** Variation in sampled function values.
 - **Smoothness:** Average second-order differences.
 - **Skewness & Kurtosis:** Statistical asymmetry and peakedness.
- **Error Estimations:**
 - **Richardson Extrapolation:** Error between coarse and fine grid integrals.
 - **PFEM Estimate:** Based on variation among particle neighborhoods.
- **Integration Method Indicator:** Binary flag for Monte Carlo vs. Quadrature.
- **Relative Error Contribution:** Region's share in total estimated error.

All features are normalized to ensure stability and generalization across diverse functions.

Action Space: 2D Integration

Continuous Action Space:

- **Action = [region_index, split_ratio, dimension, strategy]:**
 - **region_index:** Normalized index of the region to split.
 - **split_ratio:** Fraction along the chosen axis to apply the split.
 - **dimension:** 0 for vertical (x-axis) split, 1 for horizontal (y-axis).
 - **strategy:** Indicates split logic: midpoint, adaptive sample, or error-based.

Rationale:

- Enables fine-grained control over how and where to split complex 2D regions.
- Allows the agent to focus computational effort on high-error or irregular zones.
- Generalizes better to diverse shapes and function behaviors across the domain.

Reward Function: 2D Integration (Part 1)

- The reward balances accuracy, efficiency, and region-wise complexity:

$$R = \text{Global Error Reduction} + \text{Local Improvement Bonus}$$
$$+ \text{Function Behavior Rewards} + \text{Urgency} + \text{Terminal Bonus}$$

- **Global Error Term:** Change in absolute error w.r.t. true value.
- **Local Error Term:** Improvement in high-error regions post-split.
- **Function Behavior Rewards:**
 - Bonus for handling regions with high curvature, oscillation, or skew.
 - Penalizes poor splits or redundancy.

Reward Function: 2D Integration (Part 2)

- **Urgency Term:**
 - Encourages timely decisions by increasing reward pressure as the maximum allowed regions are approached.
- **Terminal Bonus:**
 - Computed at episode end, based on final global error and overall refinement quality.

Goal: Encourage accurate and efficient integration by dynamically allocating computation to complex subregions.

Architecture for 2D Adaptive Integration

Architecture Overview:

- **Algorithm:** PPO from Stable Baselines3
- **Environment:** EnhancedAdaptiveIntegrationEnv (Gym-based)
- **Input Features:**
 - Region Boundaries: (x_0, x_1, y_0, y_1)
 - Region Area, Lengths (x, y)
 - Curvature, Oscillation, Smoothness
 - Function values at center and boundaries
 - Error estimate using Richardson Extrapolation
- **Actions:**
 - Select a region to split
 - Determine split axis (x or y)
 - Adaptive split ratio based on curvature analysis
- **Reward Function:**
 - Negative of integration error over the domain
 - Penalty for excessive splits
 - Terminal bonus for efficient error reduction

Environment Logic for 2D Integration - State and Action Space

- **State Space:**

- Each region is characterized by:
 - Boundaries: (x_0, x_1, y_0, y_1)
 - Center value: $f\left(\frac{x_0+x_1}{2}, \frac{y_0+y_1}{2}\right)$
 - Curvature, Oscillation, Smoothness
 - Error estimate from Richardson Extrapolation

- **Action Space:**

- Split Axis: x or y
- Split Ratio: Adaptive based on local error analysis
- Accept current integral estimate for a region

Environment Logic for 2D Integration - Reward Structure

- **Base Reward:**

$$R_{base} = 10 \times \left(\frac{\text{Error Reduction}}{\sqrt{\text{Evaluations}}} \right)$$

- **Complexity Bonus:**

$$R_{complexity} = 2 \times \left(\frac{\text{Oscillation} + \text{Smoothness} + \text{Curvature}}{3} \right)$$

- **Gradient Reward:**

$$R_{gradient} = 3 \times \max(0, \text{Previous Gradient} - \text{Current Gradient})$$

- **Terminal Bonus:**

$$R_{terminal} = 5 \times \left(\frac{\text{Initial Error} - \text{Final Error}}{\text{Initial Error}} \right)$$

- **Total Reward:**

$$R = R_{base} + R_{complexity} + R_{gradient} + R_{terminal} - 0.1$$

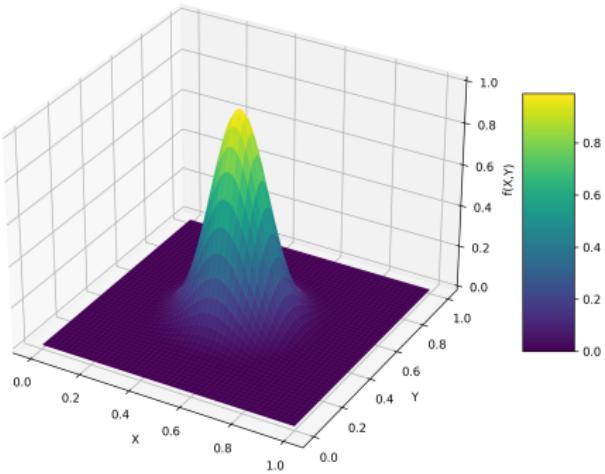
Training Details for 2D Integration

- **Training Algorithm:** PPO from Stable-Baselines3
- **Environment:** EnhancedAdaptiveIntegrationEnv
- **Training Steps:** 200,000 Timesteps
- **Learning Rate:** Linear decay from 5×10^{-4} to 1×10^{-5}
- **Training Set:**
 - Diverse function classes:
 - Highly oscillatory: $\sin(50x) \cdot \cos(50y)$
 - Singularities: $\frac{1}{\sqrt{x^2+y^2}}$
 - Discontinuous: Step and Checkerboard functions
- **Reward Shaping:**
 - Penalize unnecessary splits
 - Incentivize accurate integration with minimal evaluations

Function Classes Used in Training - Part 1

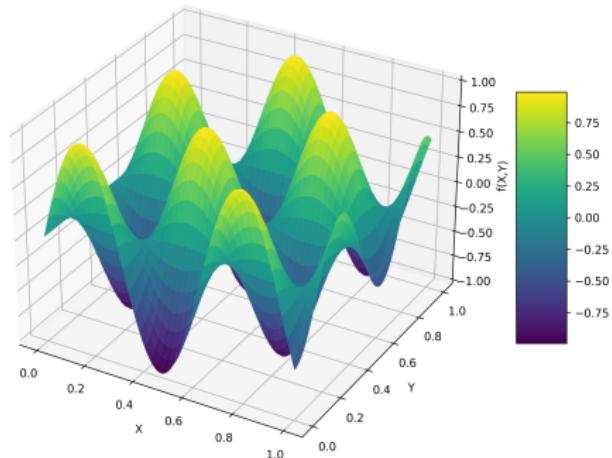
Gaussian Peak

Surface Plot: Gaussian Peak



Highly Oscillatory

Surface Plot: Highly Oscillatory

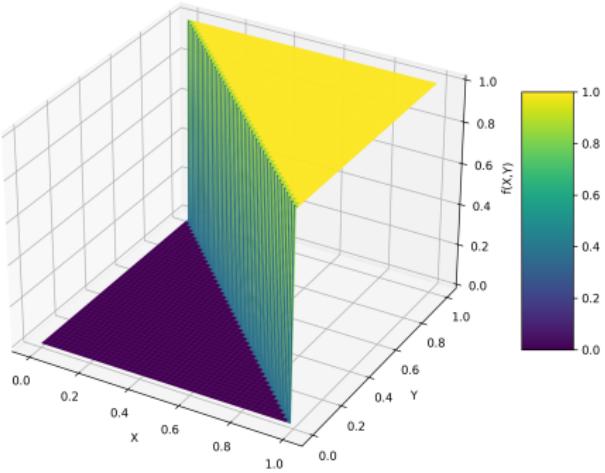


Note: The agent was trained on Gaussian peaks and oscillatory functions to capture varying function behaviors.

Function Classes Used in Training - Part 2

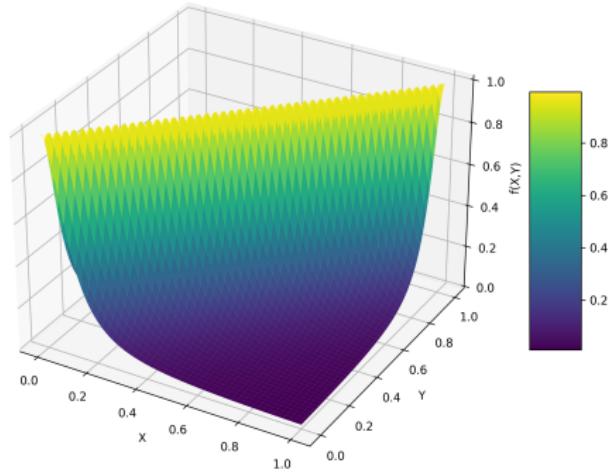
Discontinuous Graph

Surface Plot: Discontinuous Interface



Ridge Plot

Surface Plot: Ridge along $y=x$

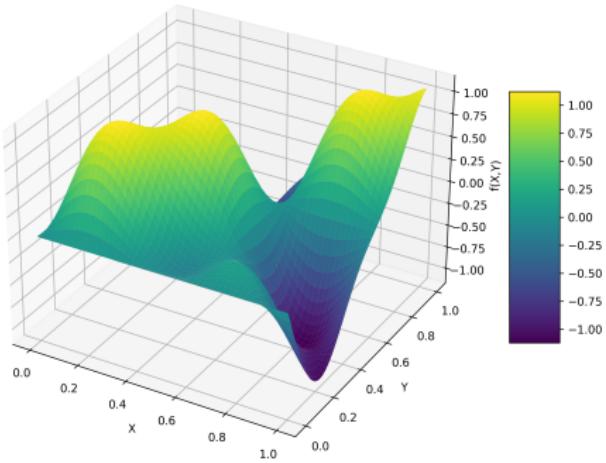


Note: The training set includes functions with sharp discontinuities and ridge structures.

Function Classes Used in Training - Part 3

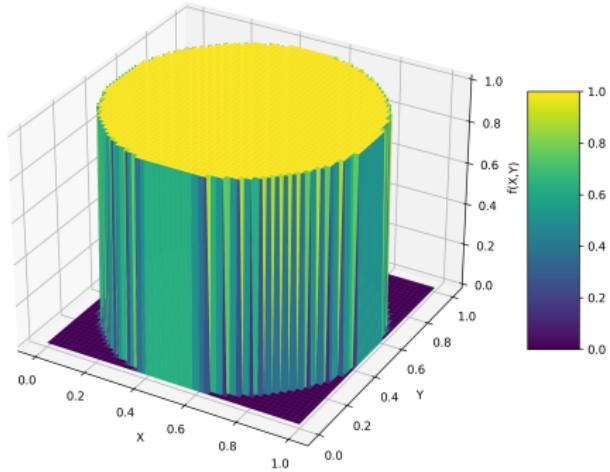
Multi Peak Function

Surface Plot: Multi-peak Function



Discontinuous Domain Function

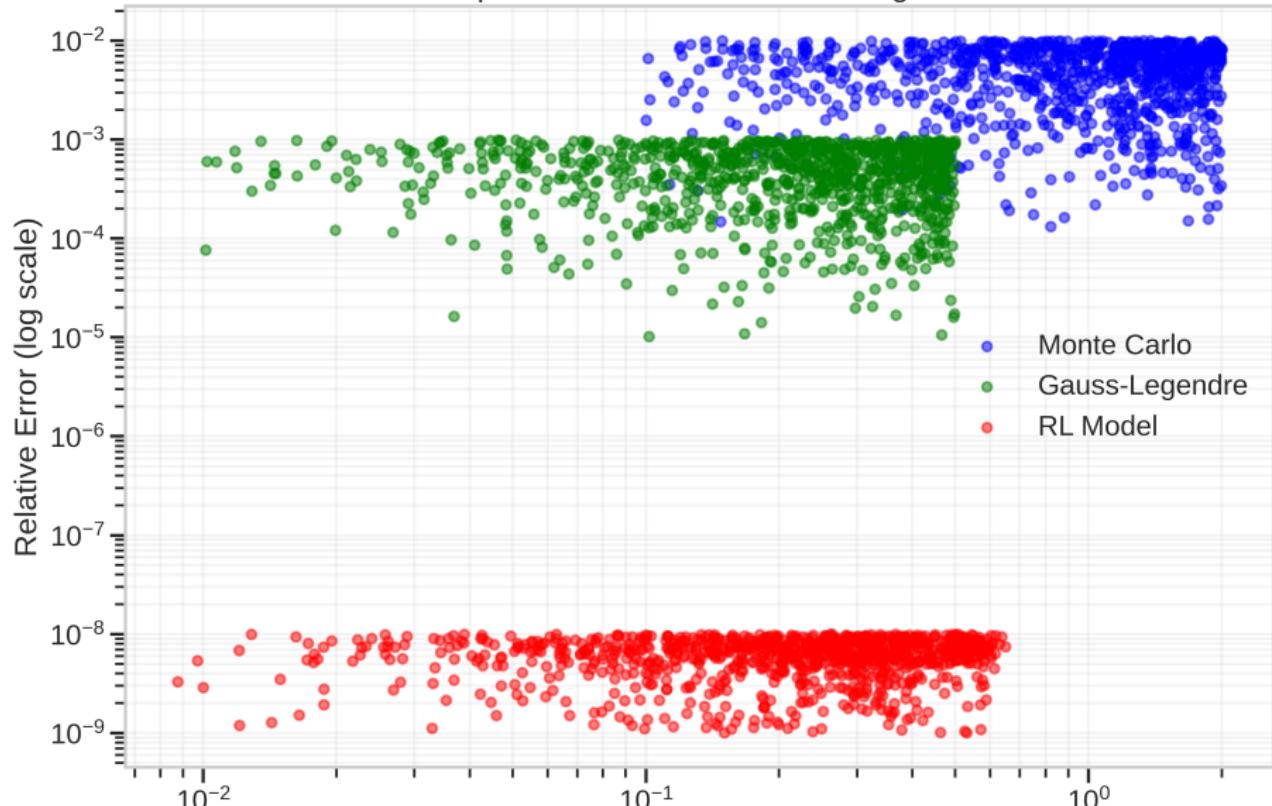
Surface Plot: Irregular Domain Indicator



Note: Multi-peak and irregular domain functions test the agent's adaptability to complex structures.

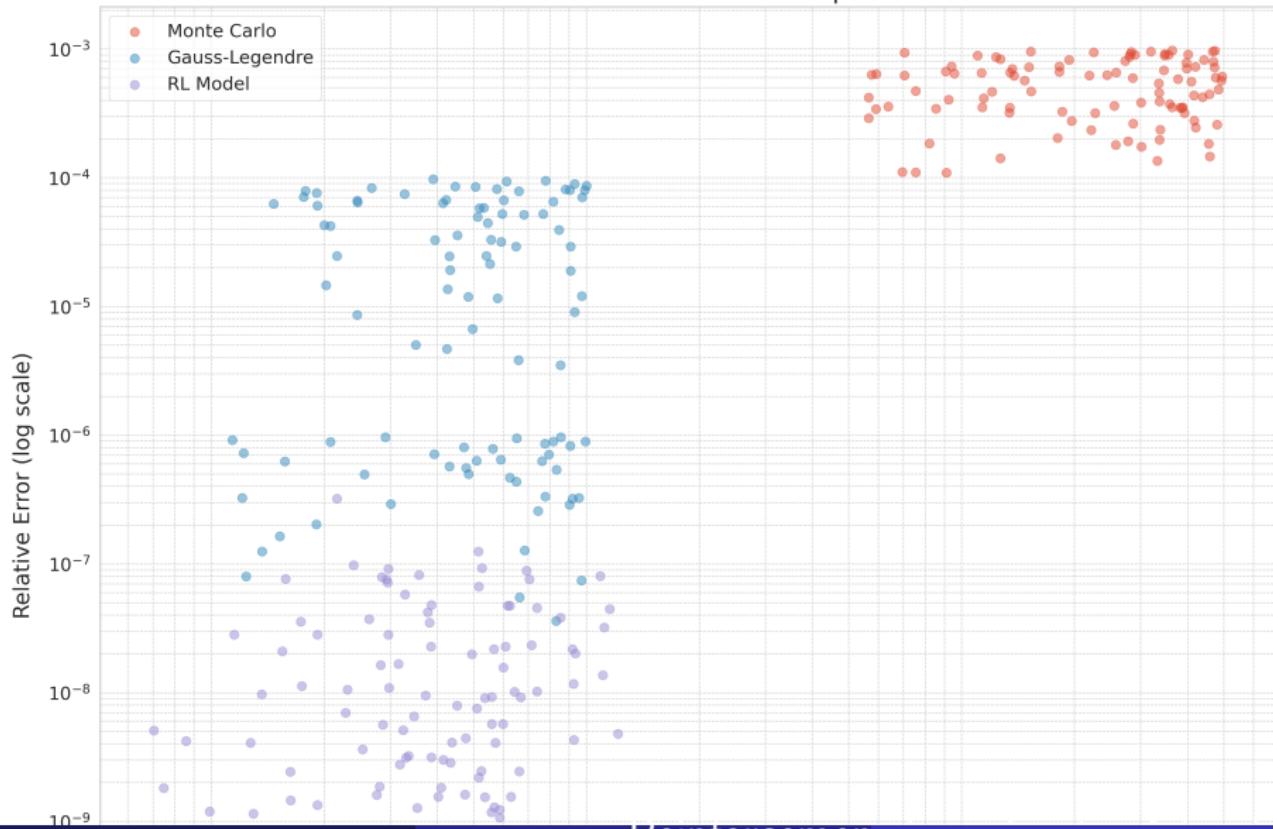
Evaluation Results

Error vs. Computation Time for Different Integration Methods



Evaluation Results

Error vs. Execution Time Comparison

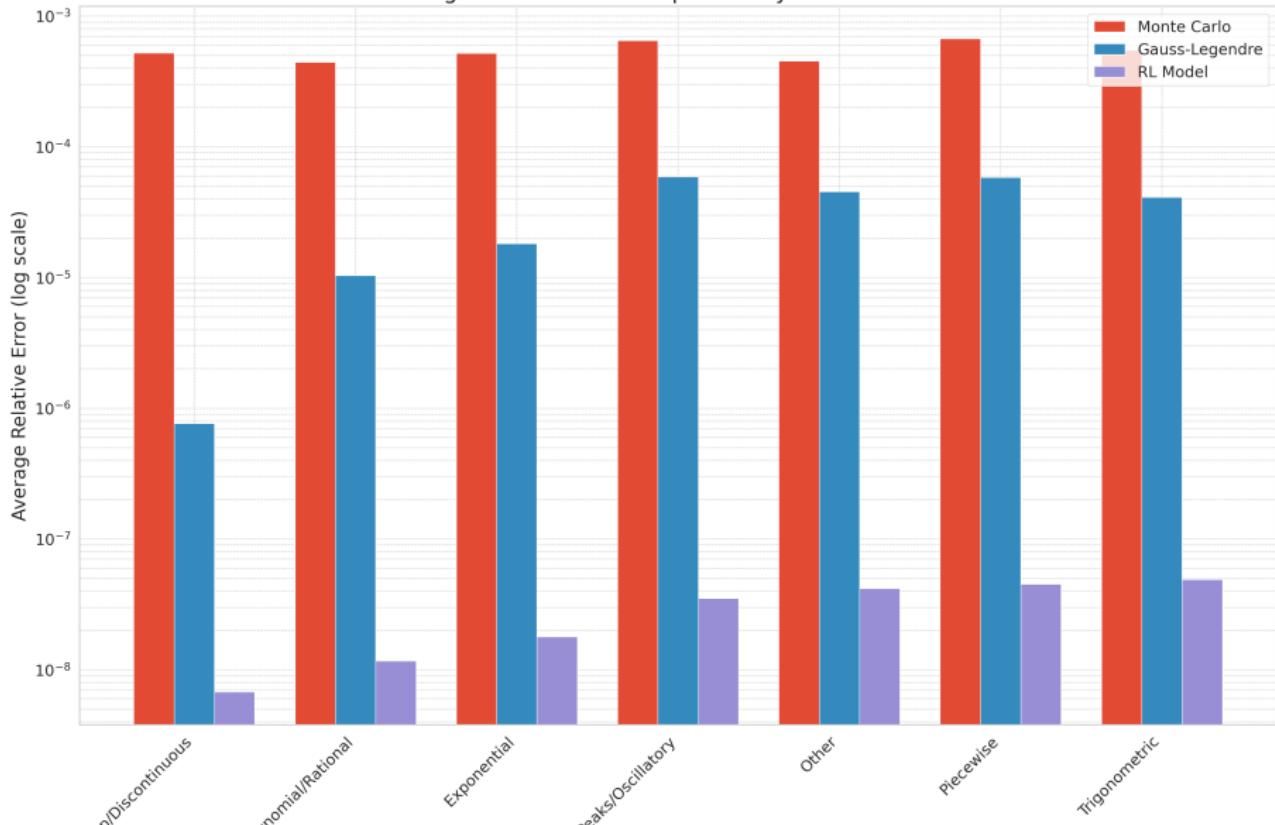


Evaluation Results



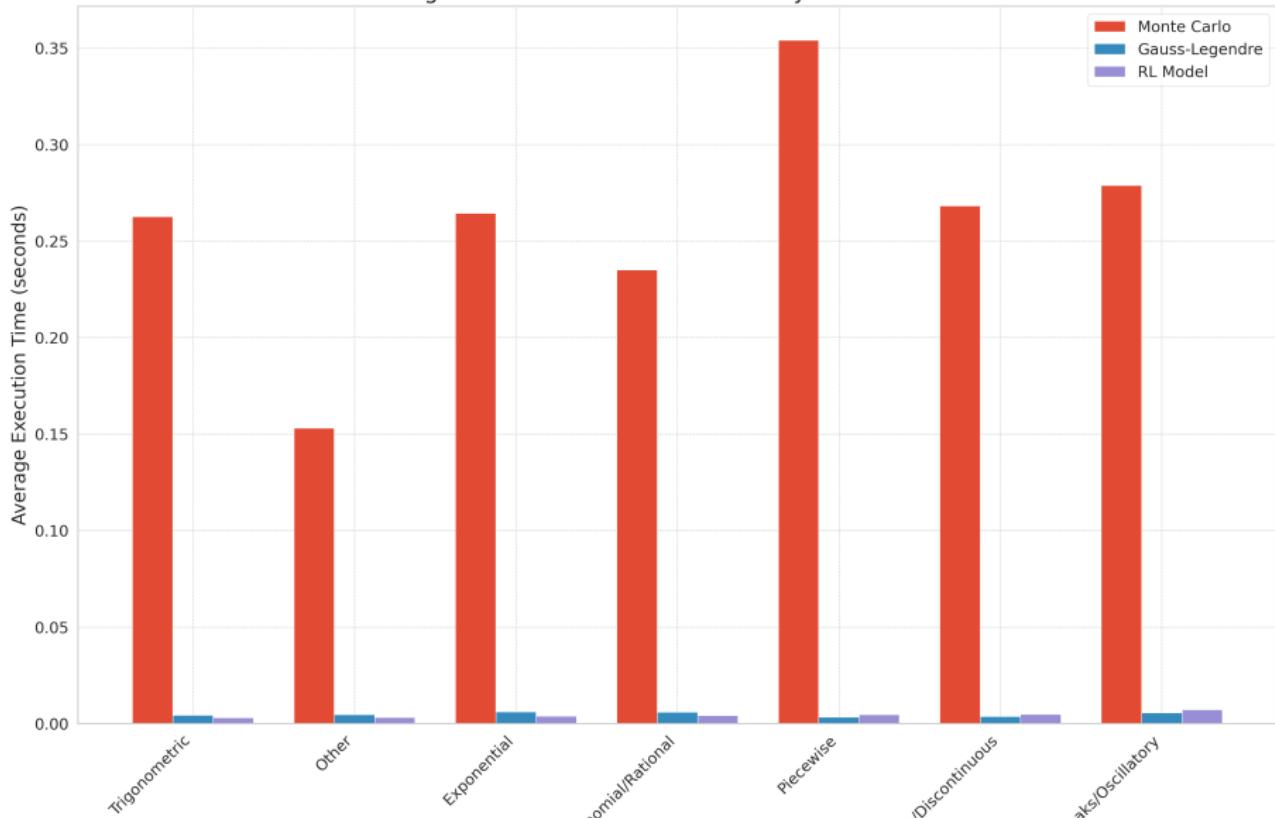
Evaluation Results

Integration Method Comparison by Function Class



Evaluation Results

Integration Method Execution Time by Function Class



Evaluation Results

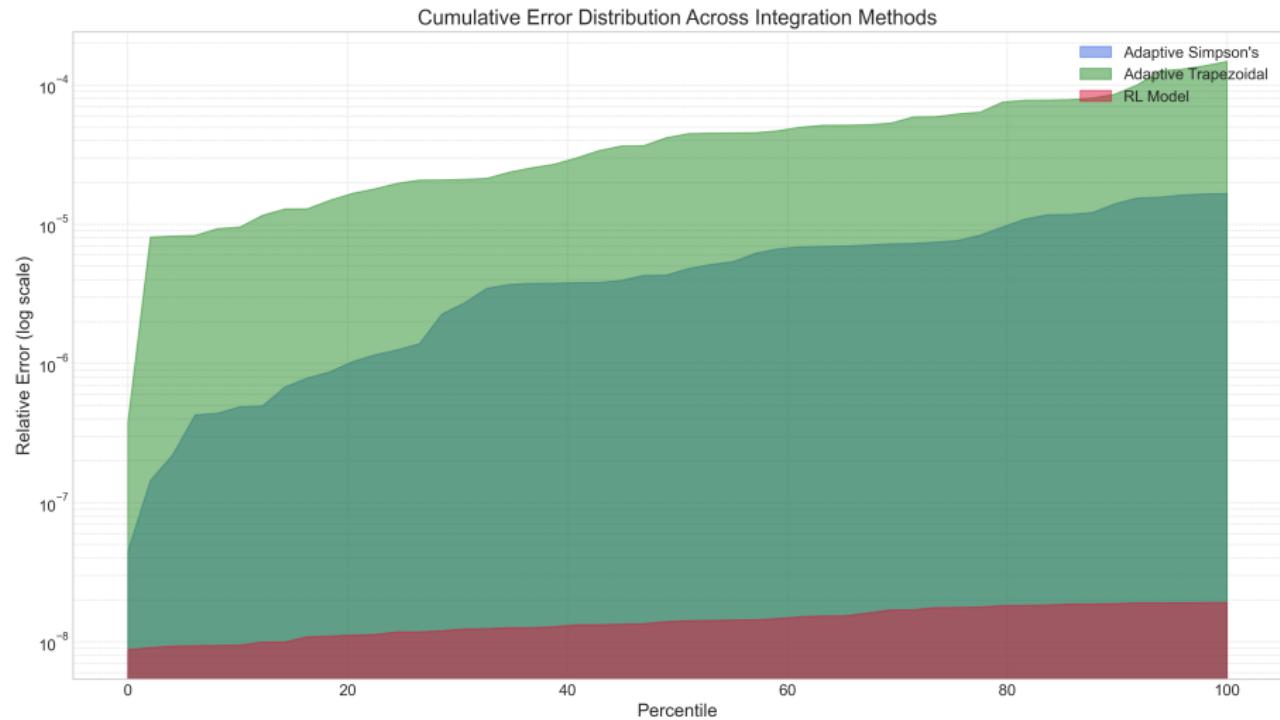


Figure: Grphah of Cummulative error across Method

Evaluation Results

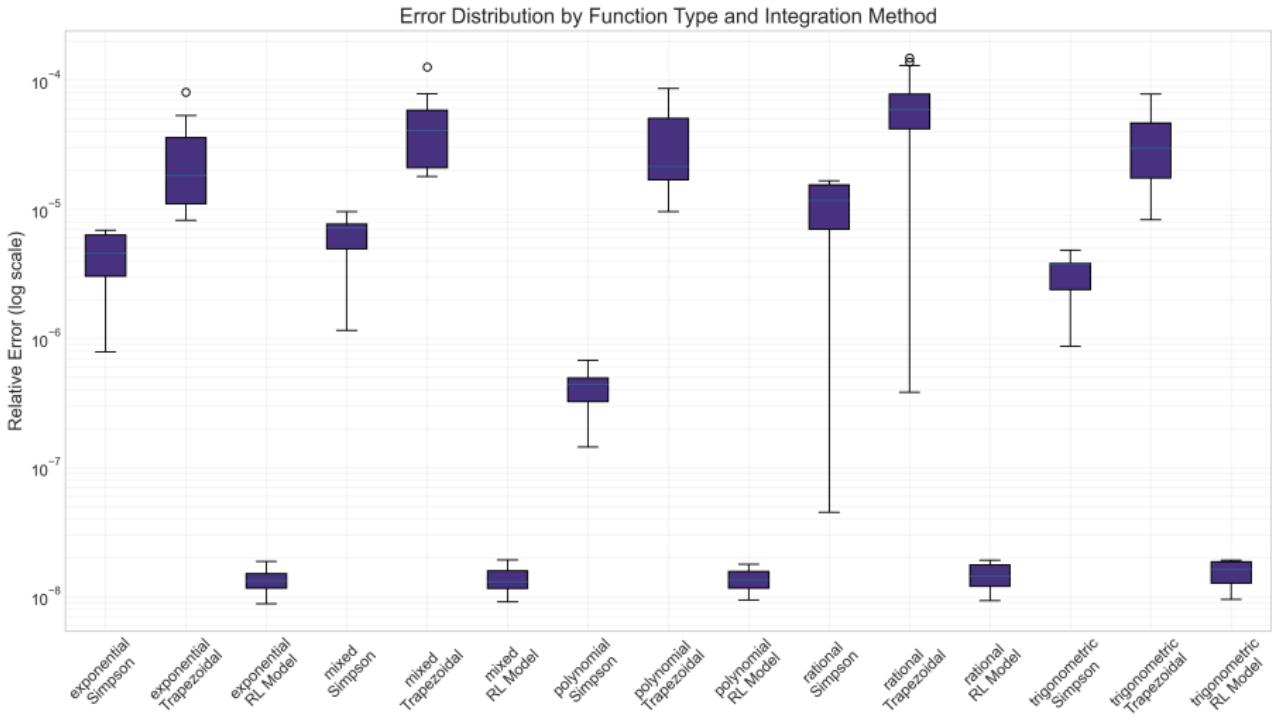


Figure: Box Plot of Error vs Function Class and Method

Evaluation Results

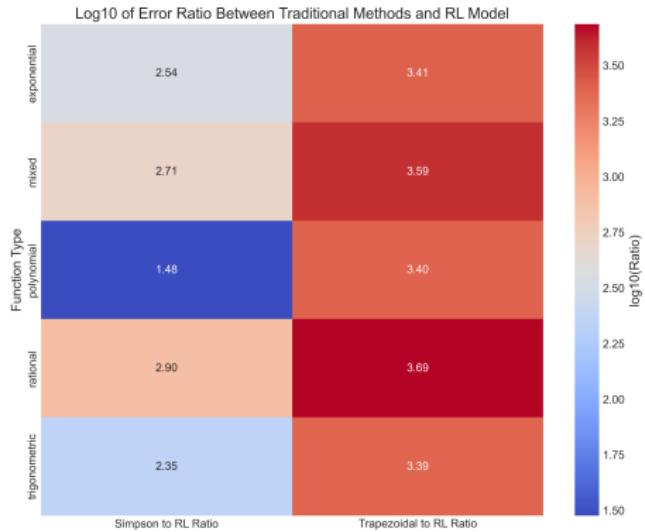
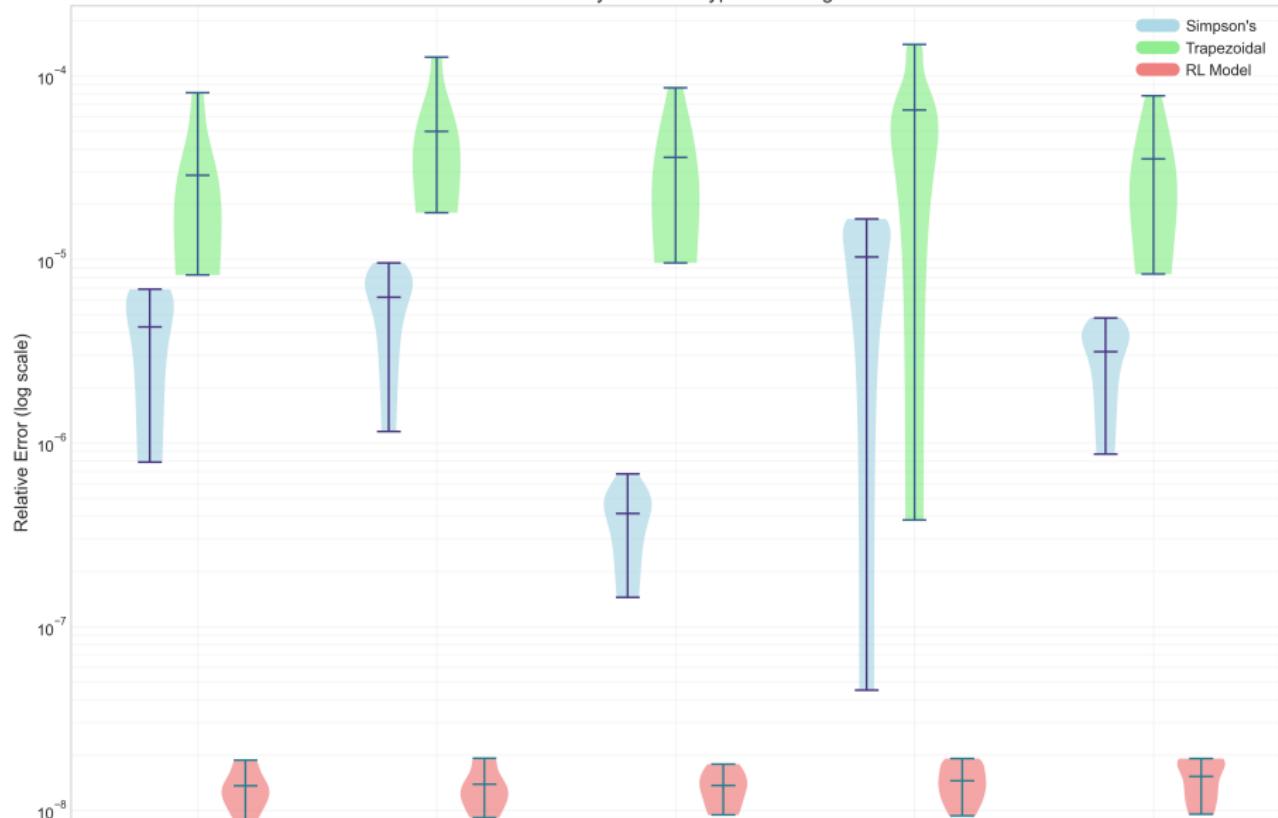


Figure: Comparison Heat Map

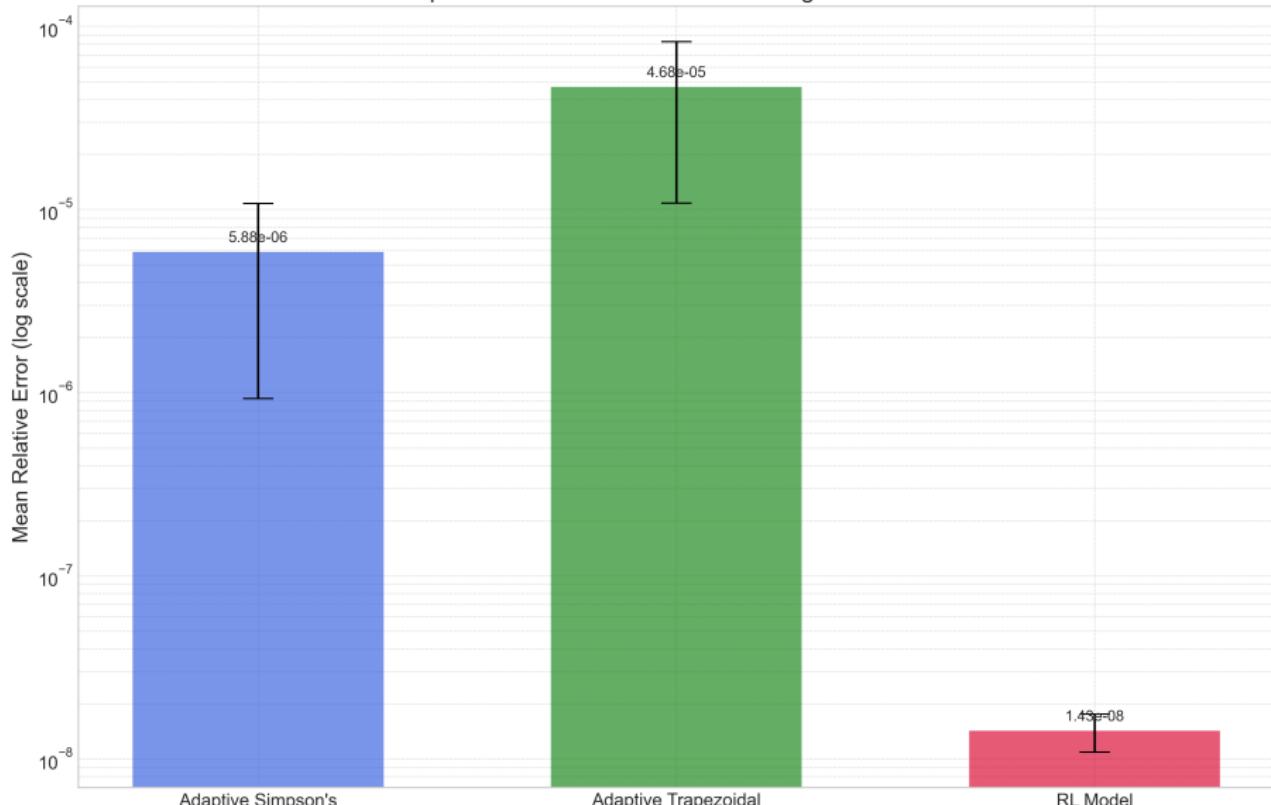
Evaluation Results

Distribution of Errors by Function Type and Integration Method



Evaluation Results

Comparison of Mean Errors Across Integration Methods



Evaluation Results



Figure: Time vs Integration Method

Evaluation Results

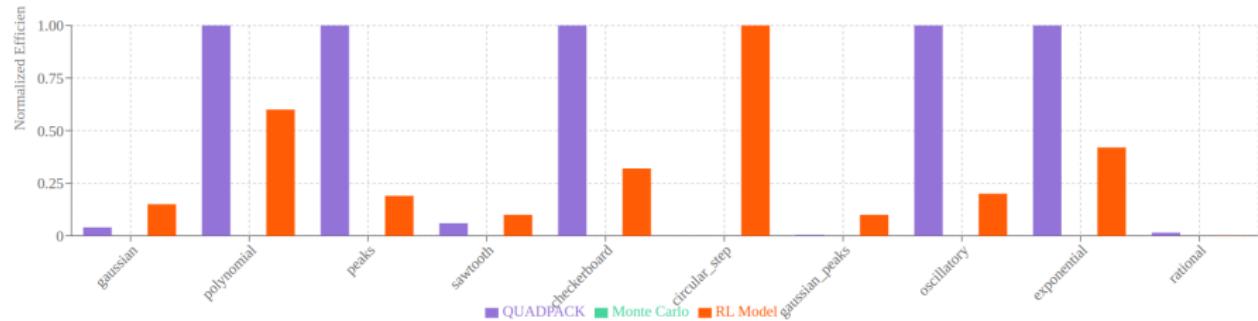


Figure: Efficiency vs Methods

Evaluation Results

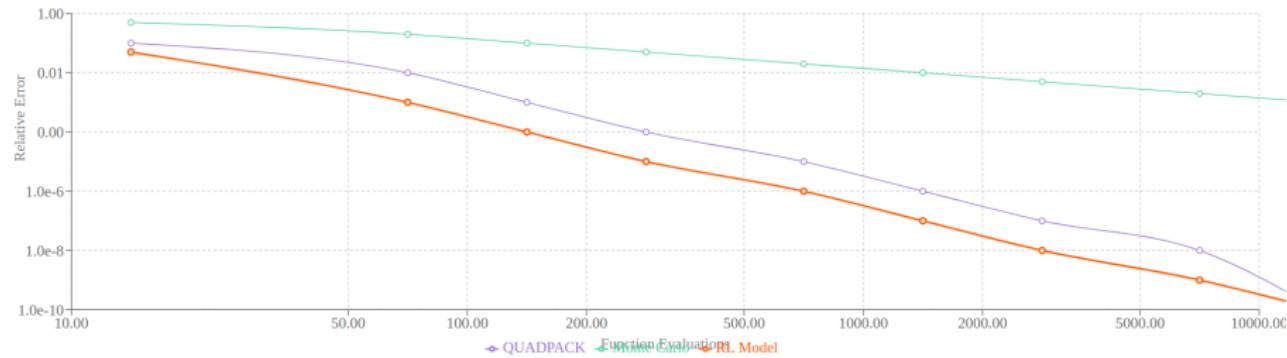
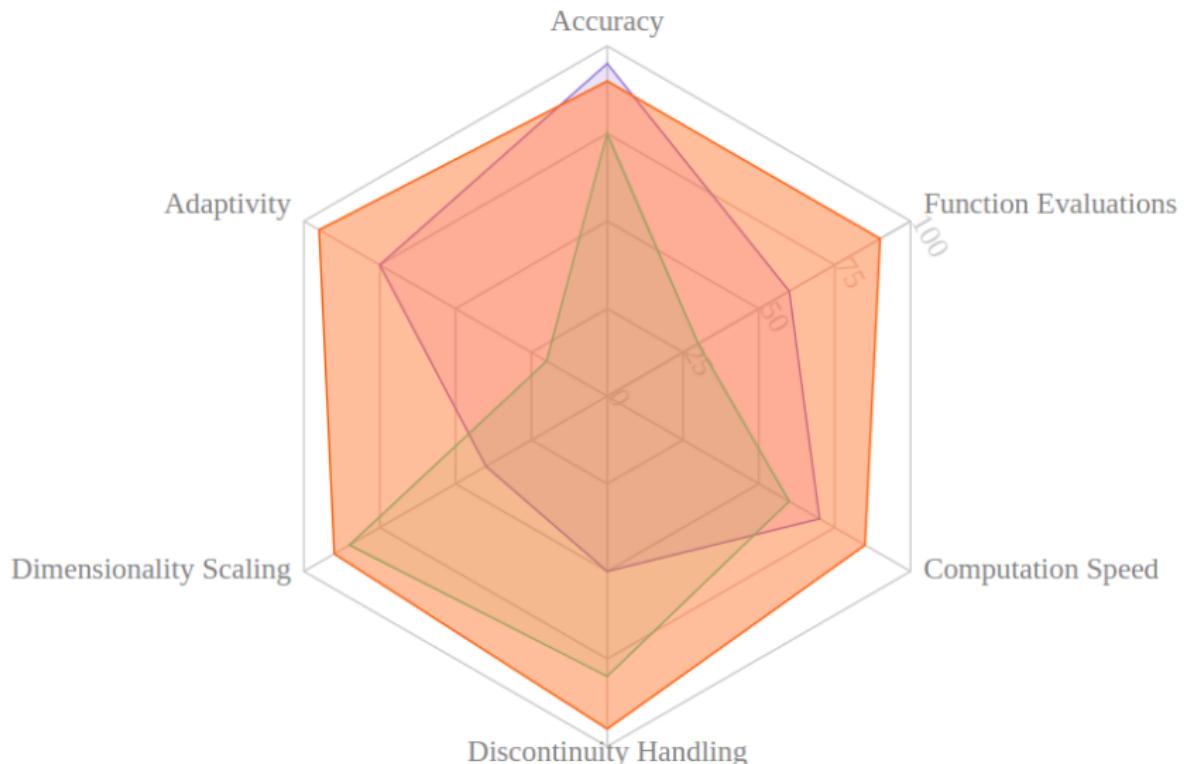


Figure: Relative Error vs Function Evaluations

Evaluation Results



Conclusion

- The implementation of reinforcement learning for adaptive 2D integration effectively allocates computational resources to regions with complex function behavior such as oscillations, discontinuities, and singularities.
- By dynamically selecting regions for refinement based on error estimates, curvature, and local function statistics, the RL agent successfully boosts the accuracy to a large extent.
- The use of multiple integration methods, including Gauss-Legendre Quadrature, Adaptive Simpson, Monte Carlo, and PFEM, enables the model to adapt to various function classes, enhancing overall robustness.

Future Work

- **Optimizing Computational Cost:**
 - Implement advanced pruning strategies to reduce unnecessary function evaluations.
 - Integrate cost-aware reward structures to prioritize high-impact regions.
- **Extending to N-Dimensional Functions:**
 - Develop scalable RL architectures capable of handling higher-dimensional integration.
 - Maintain accuracy across increasing dimensionality while managing computational complexity.

Thank You!

Questions?